

# Olympics Data Analysis

By: Purnima Agarwal

## 1) Data set used

We have taken the data set from Kaggle. It is information on the Olympic Games, from Athens 1896 to Rio 2016.

The data set has two files:

- athlete\_events.csv
- noc\_regions.csv

## 2) Understanding the data

i) We first import the needed libraries.

```
import pandas as pd
import numpy as np
import seaborn as sns
from matplotlib import pyplot as plt
```

ii) Read the files

```
ath=pd.read_csv('athlete_events.csv')
```

```
ath.head()
```

ID	Name	Sex	Age	Height	Weight	Team	NOC	Games	Year	Season	City	Sport	Event	Medal	
0	1	A Dijiang	M	24.0	180.0	80.0	China	CHN	1992 Summer	1992	Summer	Barcelona	Basketball	Basketball Men's Basketball	NaN
1	2	A Lamusi	M	23.0	170.0	60.0	China	CHN	2012 Summer	2012	Summer	London	Judo	Judo Men's Extra-Lightweight	NaN
2	3	Gunnar Nielsen Aaby	M	24.0	NaN	NaN	Denmark	DEN	1920 Summer	1920	Summer	Antwerpen	Football	Football Men's Football	NaN
3	4	Edgar Lindenau Aabye	M	34.0	NaN	NaN	Denmark/Sweden	DEN	1900 Summer	1900	Summer	Paris	Tug-Of-War	Tug-Of-War Men's Tug-Of-War	Gold
4	5	Christine Jacoba Aaftink	F	21.0	185.0	82.0	Netherlands	NED	1988 Winter	1988	Winter	Calgary	Speed Skating	Speed Skating Women's 500 metres	NaN

```
regions=pd.read_csv('noc_regions.csv')
```

```
regions.head()
```

	NOC	region	notes
0	AFG	Afghanistan	NaN
1	AHO	Curacao	Netherlands Antilles
2	ALB	Albania	NaN
3	ALG	Algeria	NaN
4	AND	Andorra	NaN

### iii) Exploring the data

- Number of rows and columns in the **athlete table**- 271116, 15 and column names in the table.

```
ath.shape
```

```
(271116, 15)
```

```
ath.columns
```

```
Index(['ID', 'Name', 'Sex', 'Age', 'Height', 'Weight', 'Team', 'NOC', 'Games',  
      'Year', 'Season', 'City', 'Sport', 'Event', 'Medal'],  
      dtype='object')
```

- Different datatypes in the table:

```
ath.dtypes
```

```
ID          int64  
Name        object  
Sex         object  
Age         float64  
Height      float64  
Weight      float64  
Team        object  
NOC          object  
Games       object  
Year        int64  
Season      object  
City        object  
Sport       object  
Event       object  
Medal       object  
dtype: object
```

- Finding the number of null values:

```
ath.isnull().sum()
```

```
ID          0
Name         0
Sex          0
Age         9474
Height      60171
Weight      62875
Team         0
NOC          0
Games        0
Year         0
Season       0
City         0
Sport        0
Event        0
Medal       231333
dtype: int64
```

- Using Describe Function:

```
ath.describe()
```

	ID	Age	Height	Weight	Year
<b>count</b>	271116.000000	261642.000000	210945.000000	208241.000000	271116.000000
<b>mean</b>	68248.954396	25.556898	175.338970	70.702393	1978.378480
<b>std</b>	39022.286345	6.393561	10.518462	14.348020	29.877632
<b>min</b>	1.000000	10.000000	127.000000	25.000000	1896.000000
<b>25%</b>	34643.000000	21.000000	168.000000	60.000000	1960.000000
<b>50%</b>	68205.000000	24.000000	175.000000	70.000000	1988.000000
<b>75%</b>	102097.250000	28.000000	183.000000	79.000000	2002.000000
<b>max</b>	135571.000000	97.000000	226.000000	214.000000	2016.000000

- Using info() function:

```
ath.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 271116 entries, 0 to 271115
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   ID          271116 non-null  int64
 1   Name        271116 non-null  object
 2   Sex         271116 non-null  object
 3   Age         261642 non-null  float64
 4   Height      210945 non-null  float64
 5   Weight      208241 non-null  float64
 6   Team        271116 non-null  object
 7   NOC         271116 non-null  object
 8   Games       271116 non-null  object
 9   Year        271116 non-null  int64
10   Season      271116 non-null  object
11   City        271116 non-null  object
12   Sport       271116 non-null  object
13   Event       271116 non-null  object
14   Medal       39783 non-null   object
dtypes: float64(3), int64(2), object(10)
memory usage: 31.0+ MB
```

- Now for the regions table

```
regions.shape
```

```
(230, 3)
```

```
regions.columns
```

```
Index(['NOC', 'region', 'notes'], dtype='object')
```

- Describing the data

```
regions.describe()
```

	NOC	region	notes
count	230	227	21
unique	230	206	21
top	ANT	Germany	Virgin Islands
freq	1	4	1

```
regions.isnull().sum()
```

```
NOC      0
region    3
notes    209
dtype: int64
```

```
regions.dtypes
```

```
NOC      object
region    object
notes    object
dtype: object
```

- Using info() function:

```
: regions.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 230 entries, 0 to 229
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0   NOC      230 non-null     object
1   region   227 non-null     object
2   notes    21 non-null      object
dtypes: object(3)
memory usage: 5.5+ KB
```

### 3) Merging the data set

We will merge both the tables for easy analysis.

We will join the two data frames using as key the NOC column with the Pandas merge() function.

```
df=ath.merge(regions,how='left',on='NOC')
```

```
df.columns
```

```
Index(['ID', 'Name', 'Sex', 'Age', 'Height', 'Weight', 'Team', 'NOC', 'Games',  
      'Year', 'Season', 'City', 'Sport', 'Event', 'Medal', 'region', 'notes'],  
      dtype='object')
```

```
df.head()
```

	ID	Name	Sex	Age	Height	Weight	Team	NOC	Games	Year	Season	City	Sport	Event	Medal	region	notes
0	1	A Dijing	M	24.0	180.0	80.0	China	CHN	1992 Summer	1992	Summer	Barcelona	Basketball	Basketball Men's Basketball	NaN	China	NaN
1	2	A Lamusi	M	23.0	170.0	60.0	China	CHN	2012 Summer	2012	Summer	London	Judo	Judo Men's Extra-Lightweight	NaN	China	NaN
2	3	Gunnar Nielsen Aaby	M	24.0	NaN	NaN	Denmark	DEN	1920 Summer	1920	Summer	Antwerpen	Football	Football Men's Football	NaN	Denmark	NaN
3	4	Edgar Lindenau Aaby	M	34.0	NaN	NaN	Denmark/Sweden	DEN	1900 Summer	1900	Summer	Paris	Tug-Of-War	Tug-Of-War Men's Tug-Of-War	Gold	Denmark	NaN

```
[16]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 271116 entries, 0 to 271115  
Data columns (total 17 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   ID           271116 non-null  int64  
1   Name         271116 non-null  object  
2   Sex          271116 non-null  object  
3   Age          261642 non-null  float64  
4   Height       210945 non-null  float64  
5   Weight       208241 non-null  float64  
6   Team         271116 non-null  object  
7   NOC          271116 non-null  object  
8   Games        271116 non-null  object  
9   Year         271116 non-null  int64  
10  Season       271116 non-null  object  
11  City         271116 non-null  object  
12  Sport        271116 non-null  object  
13  Event        271116 non-null  object  
14  Medal        39783 non-null   object  
15  region       270746 non-null  object  
16  notes        5039 non-null    object  
dtypes: float64(3), int64(2), object(12)  
memory usage: 37.2+ MB
```

## 4) Cleaning the data set

- Checking duplicate values

```
## duplicate values
```

```
df.duplicated().sum()
```

```
1385
```

This just shows there are multiple participations

```
df.isnull().sum()
```

```
ID          0
Name         0
Sex          0
Age        9474
Height     60171
Weight     62875
Team         0
NOC          0
Games       0
Year         0
Season      0
City         0
Sport        0
Event        0
Medal      231333
region       370
notes     266077
dtype: int64
```

We saw earlier that the data is mostly clean with some null values. There are null values in Age, Height, Columns, notes and Medals columns. The medal column null values are not to be removed as it represents players who haven't won a medal and notes column is not important. We can drop the other rows.

(We will use the table with dropped values in the ML part.)

- Dropping the null values:

```
df1 = df.dropna(axis=0, subset=['Age', 'Height', 'Weight', 'region'])
```

```
df1.isnull().sum()
```

```
ID          0
Name         0
Sex          0
Age          0
Height       0
Weight       0
Team         0
NOC          0
Games       0
Year         0
Season      0
City         0
Sport        0
Event        0
Medal      175723
region        0
notes     202418
dtype: int64
```

```
df1.shape
```

```
(205895, 17)
```

The number of rows is now reduced to 205895.

## 5) Analysing the data set

- Analyse for India

Taking out information about the Indian Athletes

```
df.query('Team=="India"').head(20)
```

	ID	Name	Sex	Age	Height	Weight	Team	NOC	Games	Year	Season	City	Sport	Event	Medal	region	notes
505	281	S. Abdul Hamid	M	NaN	NaN	NaN	India	IND	1928 Summer	1928	Summer	Amsterdam	Athletics	Athletics Men's 110 metres Hurdles	NaN	India	NaN
506	281	S. Abdul Hamid	M	NaN	NaN	NaN	India	IND	1928 Summer	1928	Summer	Amsterdam	Athletics	Athletics Men's 400 metres Hurdles	NaN	India	NaN
895	512	Shiny Kurisingal Abraham-Wilson	F	19.0	167.0	53.0	India	IND	1984 Summer	1984	Summer	Los Angeles	Athletics	Athletics Women's 800 metres	NaN	India	NaN
896	512	Shiny Kurisingal Abraham-Wilson	F	19.0	167.0	53.0	India	IND	1984 Summer	1984	Summer	Los Angeles	Athletics	Athletics Women's 4 x 400 metres Relay	NaN	India	NaN
897	512	Shiny Kurisingal Abraham-Wilson	F	23.0	167.0	53.0	India	IND	1988 Summer	1988	Summer	Seoul	Athletics	Athletics Women's 800 metres	NaN	India	NaN

- Countries with most medals

```
## Countries with most medals
medal_rank=df.groupby("Team")['Medal'].apply(lambda x: x.notnull().sum()).reset_index(name='Medal')
medal_rank
```

	Team	Medal
0	30. Februar	0
1	A North American Team	4
2	Acipactli	0
3	Acturus	0
4	Afghanistan	2
...	...	...
1179	Zambia	2
1180	Zefyros	0
1181	Zimbabwe	22
1182	Zut	3
1183	m-2	0

1184 rows × 2 columns



```
: medal_rank=medal_rank.sort_values("Medal",ascending=False)
medal_rank.head(10).reset_index()
```


```
:
```

	index	Team	Medal
0	1095	United States	5219
1	976	Soviet Union	2451
2	398	Germany	1984
3	412	Great Britain	1673
4	361	France	1550
5	506	Italy	1527
6	1010	Sweden	1434
7	65	Australia	1306
8	173	Canada	1243
9	476	Hungary	1127

---

- **Plotting the table using seaborn**

```
: medal=medal_rank.head(10)
sns.catplot(x="Team", y="Medal", kind="bar", data=medal,aspect=20/8.27)
: <seaborn.axisgrid.FacetGrid at 0x17e564b7520>
```



- **Teams with lowest medal count**

```
medal=medal_rank.tail(10)
medal
##sns.catplot(x="Team", y="Medal", kind="bar", data=medal,aspect=20/8.27)
```

	Team	Medal
487	India-1	0
488	India-2	0
492	Indonesia-2	0
493	Inga-LIII XXXXIII	0
494	Ingegerd	0
498	Ireland-1	0
503	Israel-1	0
504	Israel-2	0
509	Italy-3	0
1183	rn-2	0


- Total medals for India (These values count for individual players in team also)

```
### Medals India
medal_rank.loc[medal_rank['Team'] == 'India']
```

	Team	Medal
249	India	96

- Analysing for Age of the athletes

```
### Analysis of age
plt.figure(figsize=(12, 6))
plt.tight_layout()
plt.title('Age Distribution')
plt.hist(df.Age,bins=np.arange(10,80,2),color='blue',edgecolor='white')
```



Here, we see that most participants are between age 20-30.


**We also see there are athletes who are above 50. Let's see the count:**

```
df['ID'][df['Age'] > 50].count()
```

1938

- **Distribution of gold with age**

```
: ### Age distribution for gold medal
goldMedals = df[(df.Medal == 'Gold')]
goldMedals = goldMedals[np.isfinite(goldMedals['Age'])]
plt.figure(figsize=(20, 10))
plt.tight_layout()
sns.countplot(goldMedals['Age'])
plt.title('Distribution of Gold Medals')
```



- **Total teams**

```
## Number of Teams
count_team=len(pd.unique(df['Team']))
count_team
```

1184

There are 1184 total teams.


- **Medals for India information:**

```
### medals for india
india=df.loc[(df1['Team'] == 'India') & (df['Medal'].notnull())]
india
```

	ID	Name	Sex	Age	Height	Weight	Team	NOC	Games	Year	Season	City	Sport	Event	Medal	region	notes
4736	2703	Syed Mushtaq Ali	M	22.0	165.0	61.0	India	IND	1964 Summer	1964	Summer	Tokyo	Hockey	Hockey Men's Hockey	Gold	India	NaN
8192	4518	Joseph Anthony "Joe" Antic	M	29.0	168.0	59.0	India	IND	1960 Summer	1960	Summer	Roma	Hockey	Hockey Men's Hockey	Silver	India	NaN
21208	11197	Vasudevan Bhaskaran	M	29.0	174.0	68.0	India	IND	1980 Summer	1980	Summer	Moskva	Hockey	Hockey Men's Hockey	Gold	India	NaN
21815	11520	Govinda Billimogaputtaswamy	M	20.0	171.0	60.0	India	IND	1972 Summer	1972	Summer	Munich	Hockey	Hockey Men's Hockey	Bronze	India	NaN
22004	11601	Abhinav Bindra	M	25.0	173.0	70.0	India	IND	2008 Summer	2008	Summer	Beijing	Shooting	Shooting Men's Air Rifle, 10 metres	Gold	India	NaN


- **Plotting for Gold, Silver, Bronze:**

```
: sns.countplot(india['Medal'])
```



- **Gender Distribution**

```
plt.figure(figsize=(12, 6))  
plt.tight_layout()  
plt.title('Gender Distribution')  
plt.pie(gender_count, labels=gender_count.index, autopct='%1.1f%%')
```



- Medal won by female participants

```

: ### Medal won by female Participants
female_medal=df[(df.Sex=='F')]
female_medal.Medal.value_counts()

: Bronze      3771
  Gold        3747
  Silver      3735
  Name: Medal, dtype: int64

```


- Female participation in each Summer Olympics

```

female=df[(df.Sex=='F') & (df.Season=='Summer')][['Sex', 'Year']]
female=female.groupby('Year').count().reset_index()
female.tail()

```

	Year	Sex
23	2000	5431
24	2004	5546
25	2008	5816
26	2012	5815
27	2016	6223




In the graph we can see that woman participation has increased in the recent years.

- Correlation among the data
- Plotting the heatmap

```
c=win.corr()
c
```

	Age	Weight	Height	Medal
Age	1.000000	0.211951	0.141736	0.028036
Weight	0.211951	1.000000	0.796652	0.088426
Height	0.141736	0.796652	1.000000	0.089117
Medal	0.028036	0.088426	0.089117	1.000000

```
plt.figure(figsize=(12,6))
sns.heatmap(c,annot=True)
plt.show()
```



## Machine learning model to Predict winning

- Creating a new data frame

## 1-Winning a medal

## 0- Losing

```
### Taking out the numerical Values using the table with non null values
win=df1[['Age','Weight','Height','Medal']]
win.loc[df1['Medal'].notnull(), 'Medal'] = 1
win.loc[df1['Medal'].isnull(), 'Medal'] = 0
win
```

	Age	Weight	Height	Medal
0	24.0	80.0	180.0	0
1	23.0	60.0	170.0	0
4	21.0	82.0	185.0	0
5	21.0	82.0	185.0	0
6	25.0	82.0	185.0	0
...	...	...	...	...
271111	29.0	89.0	179.0	0
271112	27.0	59.0	176.0	0
271113	27.0	59.0	176.0	0
271114	30.0	96.0	185.0	0
271115	34.0	96.0	185.0	0

205895 rows × 4 columns

```
win['Medal'].value_counts()
```

```
0    175723
1     30172
Name: Medal, dtype: int64
```

```
import pylab as pl
import numpy as np
import scipy.optimize as opt
from sklearn import preprocessing
%matplotlib inline
```

```
X = np.asarray(win[['Age', 'Weight', 'Height']]).astype('int')
X[0:5]
```

```
array([[ 24,  80, 180],
       [ 23,  60, 170],
       [ 21,  82, 185],
       [ 21,  82, 185],
       [ 25,  82, 185]])
```

```
y = np.asarray(win['Medal']).astype('int')
y[0:5]
```

```
array([0, 0, 0, 0, 0])
```

```
from sklearn import preprocessing
X = preprocessing.StandardScaler().fit(X).transform(X)
X[0:5]
```

```
array([[ -0.19,  0.65,  0.44],
       [-0.38, -0.75, -0.51],
```

## normalizing the values

- **Logistic Regression**



```

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
X_train, X_test, y_train, y_test = train_test_split( X, y, test_size=0.2, random_state=4)
LR = LogisticRegression(C=0.01, solver='liblinear').fit(X_train,y_train)
yhat = LR.predict(X_test)

```

yhat

```
array([0, 0, 0, ..., 0, 0, 0])
```

```

yhat_prob = LR.predict_proba(X_test)
yhat_prob

```

```

array([[0.85, 0.15],
       [0.83, 0.17],
       [0.85, 0.15],
       ...,
       [0.82, 0.18],
       [0.84, 0.16],
       [0.86, 0.14]])

```

- Let's try Root mean Squared error for error calculation

```

## root mean squared error
from sklearn.metrics import mean_squared_error

np.sqrt(mean_squared_error(y_test, yhat))

```

```
0.3851653093161678
```

- Plotting confusion matrix

```

from sklearn.metrics import classification_report, confusion_matrix
import itertools
def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

```

```
plt.imshow(cm, interpolation='nearest', cmap=cmap)
plt.title(title)
plt.colorbar()
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=45)
plt.yticks(tick_marks, classes)

fmt = '.2f' if normalize else 'd'
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], fmt),
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')
print(confusion_matrix(y_test, yhat, labels=[1,0]))
```

```
[[ 0 6109]
 [ 0 35070]]
```


We are plotting a confusion matrix to evaluate the performance.

```
cnf_matrix = confusion_matrix(y_test, yhat, labels=[1,0])
np.set_printoptions(precision=2)

# Plot non-normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=['Medal Win=1','Medal Loss=0'],normalize= False,  title='Confusion matrix')
```

Confusion matrix, without normalization

```
[[ 0 6109]
 [ 0 35070]]
```



Based on the count of each section, we can calculate precision and recall of each label:

- **Precision** is a measure of the accuracy provided that a class label has been predicted.
- **Recall** is the true positive rate.

So, we can calculate the precision and recall of each class.

**F1 score:** Now we are in the position to calculate the F1 scores for each label based on the precision and recall of that label.

The F1 score is the harmonic average of the precision and recall, where an F1 score reaches its best value at 1 (perfect precision and recall) and worst at 0. It is a good way to show that a classifier has a good value for both recall and precision.

```
print(classification_report(y_test, yhat))
```

	precision	recall	f1-score	support
0	0.85	1.00	0.92	35070
1	0.00	0.00	0.00	6109
accuracy			0.85	41179
macro avg	0.43	0.50	0.46	41179
weighted avg	0.73	0.85	0.78	41179

**# We see here accuracy of the model is 0.85**

- **Decision Tree Classification**

```
from sklearn.tree import DecisionTreeClassifier
```

```
## We will first create an instance of the DecisionTreeClassifier called Tree
Tree = DecisionTreeClassifier(criterion="entropy", max_depth = 4)
Tree.fit(X_train,y_train)
## predict the value
predTree =Tree.predict(X_test)
```

```
predTree[0:10]
```

```
array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0])
```

```
## defining the accuracy of the model|
from sklearn import metrics
print("DecisionTrees's Accuracy: ", metrics.accuracy_score(y_test, predTree))
```

```
DecisionTrees's Accuracy:  0.8516476844993808
```

- Let's try Root mean Squared error for error calculation

```
np.sqrt(mean_squared_error(y_test, predTree))
```

0.3851653093161678

```
cnf_matrix = confusion_matrix(y_test, predTree, labels=[1,0])  
np.set_printoptions(precision=2)
```

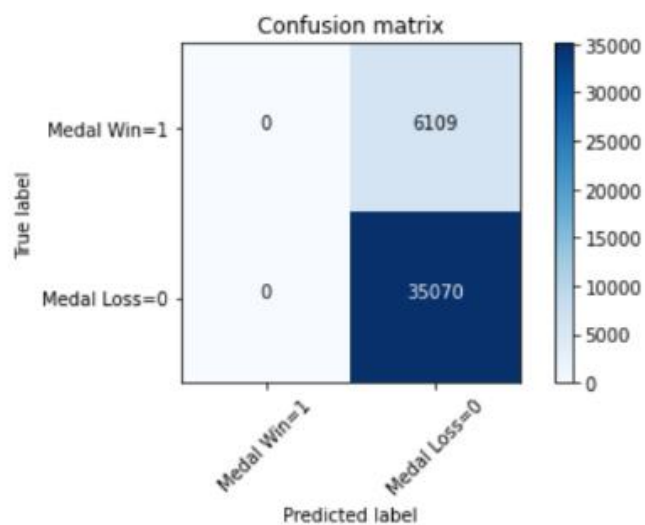
```
# Plot non-normalized confusion matrix
```

```
plt.figure()
```

```
plot_confusion_matrix(cnf_matrix, classes=['Medal Win=1','Medal Loss=0'],normalize= False,  title='Confusion matrix')
```

Confusion matrix, without normalization

```
[[ 0 6109]  
 [ 0 35070]]
```



```
print (classification_report(y_test, predTree))
```

	precision	recall	f1-score	support
0	0.85	1.00	0.92	35070
1	0.00	0.00	0.00	6109
accuracy			0.85	41179
macro avg	0.43	0.50	0.46	41179
weighted avg	0.73	0.85	0.78	41179

**# We see here accuracy of the model is 0.85**

- **K-Nearest Neighbour**

Import the files and predict the values and then calculate root mean squared error

```
### k nearest neighbour
```

```
from sklearn.neighbors import KNeighborsClassifier
k = 4
#Train Model and Predict
neigh = KNeighborsClassifier(n_neighbors = k).fit(X_train,y_train)
yhat = neigh.predict(X_test)
from sklearn import metrics
print("Train set Accuracy: ", metrics.accuracy_score(y_train, neigh.predict(X_train)))
print("Test set Accuracy: ", metrics.accuracy_score(y_test, yhat))
```

```
Train set Accuracy: 0.8548896282085529
Test set Accuracy: 0.8427596590495156
```

```
##RMSE
```

```
np.sqrt(mean_squared_error(y_test, yhat))
```

```
0.39343059676883585
```

- **Confusion matrix**

```
cnf_matrix = confusion_matrix(y_test, yhat, labels=[1,0])
np.set_printoptions(precision=2)
```


```
# Plot non-normalized confusion matrix
```

```
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=['Medal Win=1','Medal Loss=0'],normalize= False,  title='Confusion matrix')
```

Confusion matrix, without normalization

```
[[ 147 5962]
 [ 412 34658]]
```

--



- **Calculate the accuracy**

# We see here accuracy of the model is 0.85

```

: print (classification_report(y_test, yhat))

```

	precision	recall	f1-score	support
0	0.85	0.99	0.92	35070
1	0.26	0.02	0.04	6109
accuracy			0.85	41179
macro avg	0.56	0.51	0.48	41179
weighted avg	0.77	0.85	0.79	41179

K in KNN, is the number of nearest neighbours to examine. It is supposed to be specified by the user. So, how can we choose right value for K? The general solution is to reserve a part of your data for testing the accuracy of the model. Then choose  $k=1$ , use the training part for modelling, and calculate the accuracy of prediction using all samples in your test set. Repeat this process, increasing the  $k$ , and see which  $k$  is the best for your model.

We can calculate the accuracy of KNN for different values of  $k$ .

```

## We can calculate the accuracy of KNN for different values of k.

```

```

Ks=10
mean_acc = np.zeros((Ks-1))
std_acc = np.zeros((Ks-1))


for n in range(1,Ks):

    #Train Model and Predict
    neigh = KNeighborsClassifier(n_neighbors = n).fit(X_train,y_train)
    yhat=neigh.predict(X_test)
    mean_acc[n-1] = metrics.accuracy_score(y_test, yhat)

    std_acc[n-1]=np.std(yhat==y_test)/np.sqrt(yhat.shape[0])

mean_acc
plt.plot(range(1,Ks),mean_acc,'g')
plt.fill_between(range(1,Ks),mean_acc - 1 * std_acc,mean_acc + 1 * std_acc, alpha=0.10)
plt.fill_between(range(1,Ks),mean_acc - 3 * std_acc,mean_acc + 3 * std_acc, alpha=0.10,color="green")
plt.legend(('Accuracy ', '+/- 1xstd', '+/- 3xstd'))
plt.ylabel('Accuracy ')
plt.xlabel('Number of Neighbors (K)')
plt.tight_layout()
plt.show()

```



## Accuracy:

- Logistic Regression: 0.85
- Decision Tree: 0.85
- K- Nearest Neighbour: 0.85

## Conclusion:

- USA has the most medals
- We see that most participants are between age 20-30.
- Female participation has increased over the years.
- Most gold is won by participants that are between age 20-30.
- People above age 50 are also winning gold medal
- Winning doesn't depend much on age, weight or height.
- We see here that the accuracy shown by all the three models are equal as winning a medal depends on many factors and we have insufficient data for that.
- We need more data to accurately predict the win.